



Chapter 3: Operating-System Structures

- System Components
- Operating System Services
- System Calls
- System Programs
- System Structure
- Virtual Machines
- System Design and Implementation
- System Generation





Common System Components

- Process Management
- Main Memory Management
- File Management
- I/O System Management
- Secondary Management
- Networking
- Protection System
- Command-Interpreter System





Process Management

- A *process* is a program in execution
 - A process needs certain resources, including CPU time, memory, files, and I/O devices, to accomplish its task
- The operating system is responsible for the following activities in connection with process management
 - Process creation and deletion
 - Process suspension and resumption
 - Provision of mechanisms for:
 - ▶ process synchronization
 - ▶ process communication





Main-Memory Management

- *Memory* is a large array of words or bytes, each with its own address
 - It is a repository of quickly accessible data shared by the CPU and I/O devices
- *Main memory* is a volatile storage device. It loses its contents in the case of system failure
- The operating system is responsible for the following activities in connections with memory management
 - Keep track of which parts of memory are currently being used and by whom
 - Decide which processes to load when memory space becomes available
 - Allocate and deallocate memory space as needed





File Management

- A *file* is a collection of related information defined by its creator
 - Commonly, files represent programs (both source and object forms) and data
- The operating system is responsible for the following activities in connections with file management:
 - File creation and deletion
 - Directory creation and deletion
 - Support of primitives for manipulating files and directories
 - Mapping files onto *secondary storage*
 - File backup on *stable (nonvolatile) storage* media





I/O System Management

- The I/O system consists of:
 - A buffer-caching system
 - A general device-driver interface
 - Drivers for specific hardware devices





Secondary-Storage Management

- Since main memory (*primary storage*) is volatile and too small to accommodate all data and programs permanently, the computer system must provide *secondary storage* to back up main memory
- Most modern computer systems use disks as the principle on-line storage medium, for both programs and data
- The operating system is responsible for the following activities in connection with disk management:
 - Free space management
 - Storage allocation
 - Disk scheduling





Networking (Distributed Systems)

- A *distributed* system is a collection processors that do not share memory or a clock
 - Each processor has its own local memory
- The processors in the system are connected through a communication network
- Communication takes place using a *protocol*
- A distributed system provides user access to various system resources
- Access to a shared resource allows:
 - Computation speed-up
 - Increased data availability
 - Enhanced reliability





Protection System

- *Protection* refers to a mechanism for controlling access by programs, processes, or users to both system and user resources
- The protection mechanism must:
 - distinguish between authorized and unauthorized usage
 - specify the controls to be imposed
 - provide a means of enforcement





Command-Interpreter System

- Many commands are given to the operating system by control statements which deal with:
 - Process creation and management
 - I/O handling
 - Secondary-storage management
 - Main-memory management
 - File-system access
 - Protection
 - Networking





Command-Interpreter System (Cont.)

- The program that reads and interprets control statements is called variously:
 - *command-line interpreter*
 - *shell (in UNIX)*

Its function is to get and execute the next command statement





Operating System Services

- Program execution – system capability to load a program into memory and to run it
- I/O operations – since user programs cannot execute I/O operations directly, the operating system must provide some means to perform I/O
- File-system manipulation – program capability to read, write, create, and delete files
- Communications – exchange of information between processes executing either on the same computer or on different systems tied together by a network. Implemented via *shared memory* or *message passing*
- Error detection – ensure correct computing by detecting errors in the CPU and memory hardware, in I/O devices, or in user programs





Additional Operating System Functions

Additional functions exist not for helping the user, but rather for ensuring efficient system operations

- Resource allocation – allocating resources to multiple users or multiple jobs running at the same time
- Accounting – keep track of and record which users use how much and what kinds of computer resources for account billing or for accumulating usage statistics
- Protection – ensuring that all access to system resources is controlled





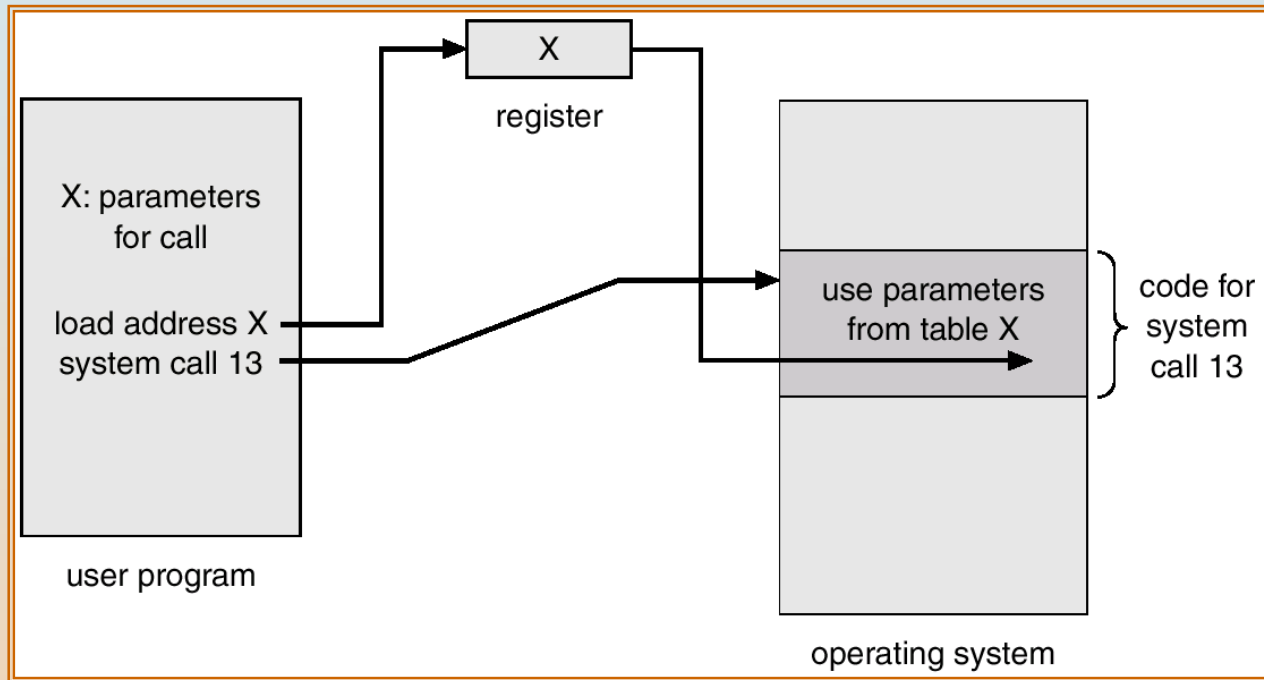
System Calls

- *System calls* provide the interface between a running program and the operating system
 - Generally available as assembly-language instructions
 - Languages defined to replace assembly language for systems programming allow system calls to be made directly (e.g., C, C++)
- Three general methods are used to *pass parameters* between a running program and the operating system
 - Pass parameters in *registers*
 - Store the parameters in a table in memory, and the table address is passed as a parameter in a register
 - *Push* (store) the parameters onto the *stack* by the program, and *pop* off the stack by operating system





Passing of Parameters As A Table





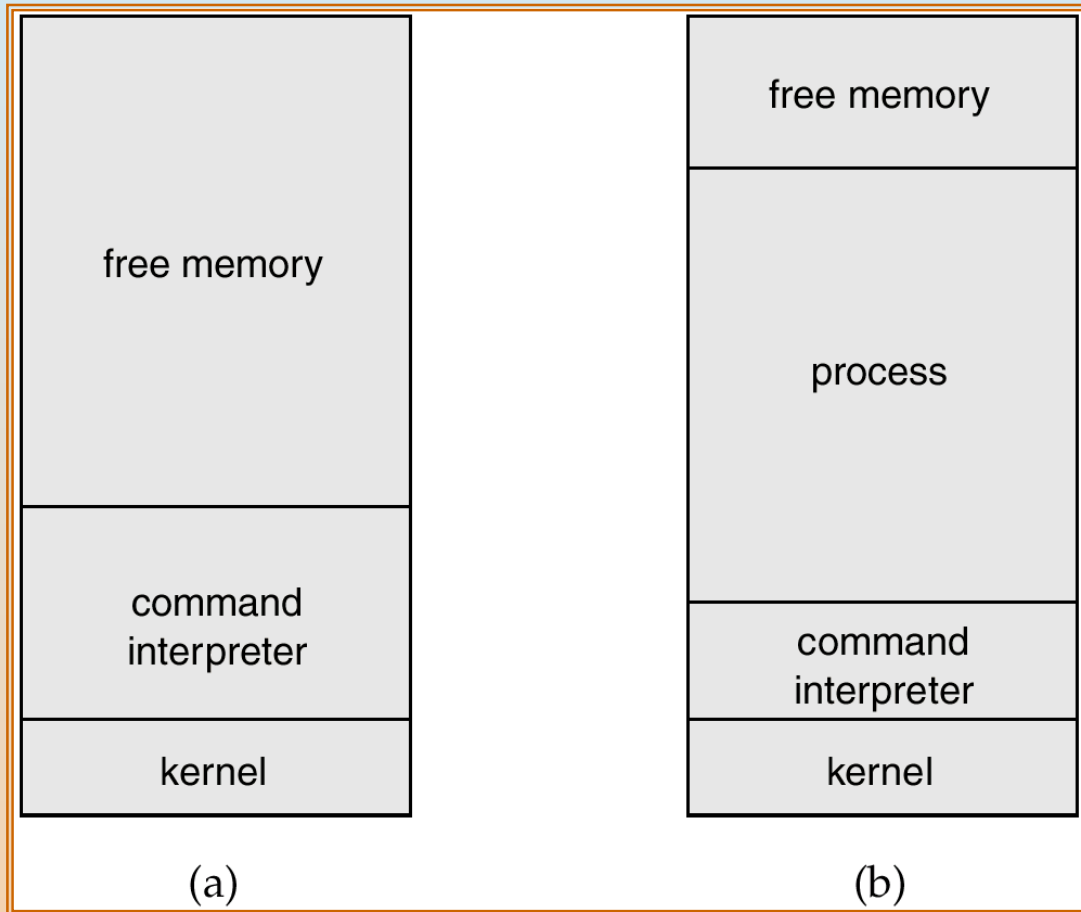
Types of System Calls

- Process control
- File management
- Device management
- Information maintenance
- Communications





MS-DOS Execution



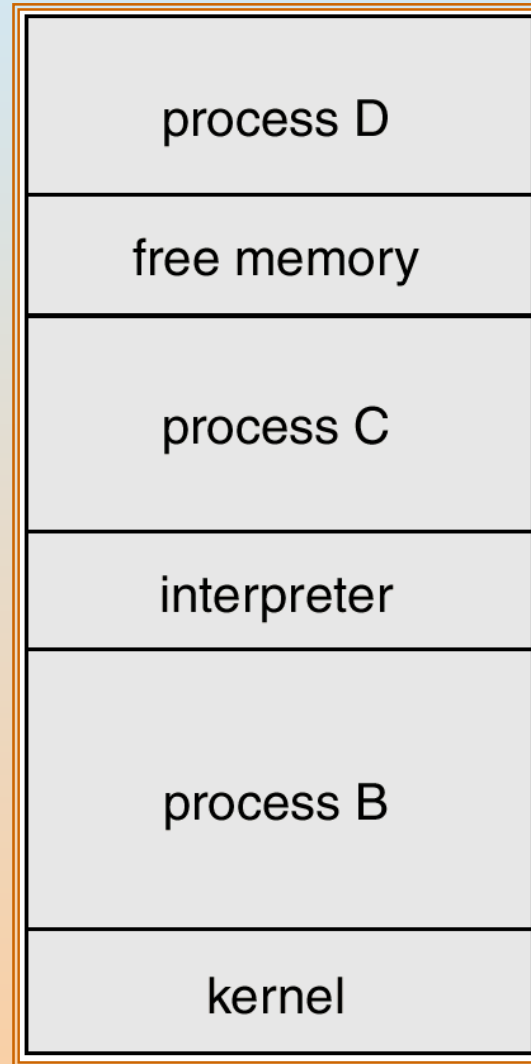
At System Start-up

Running a Program





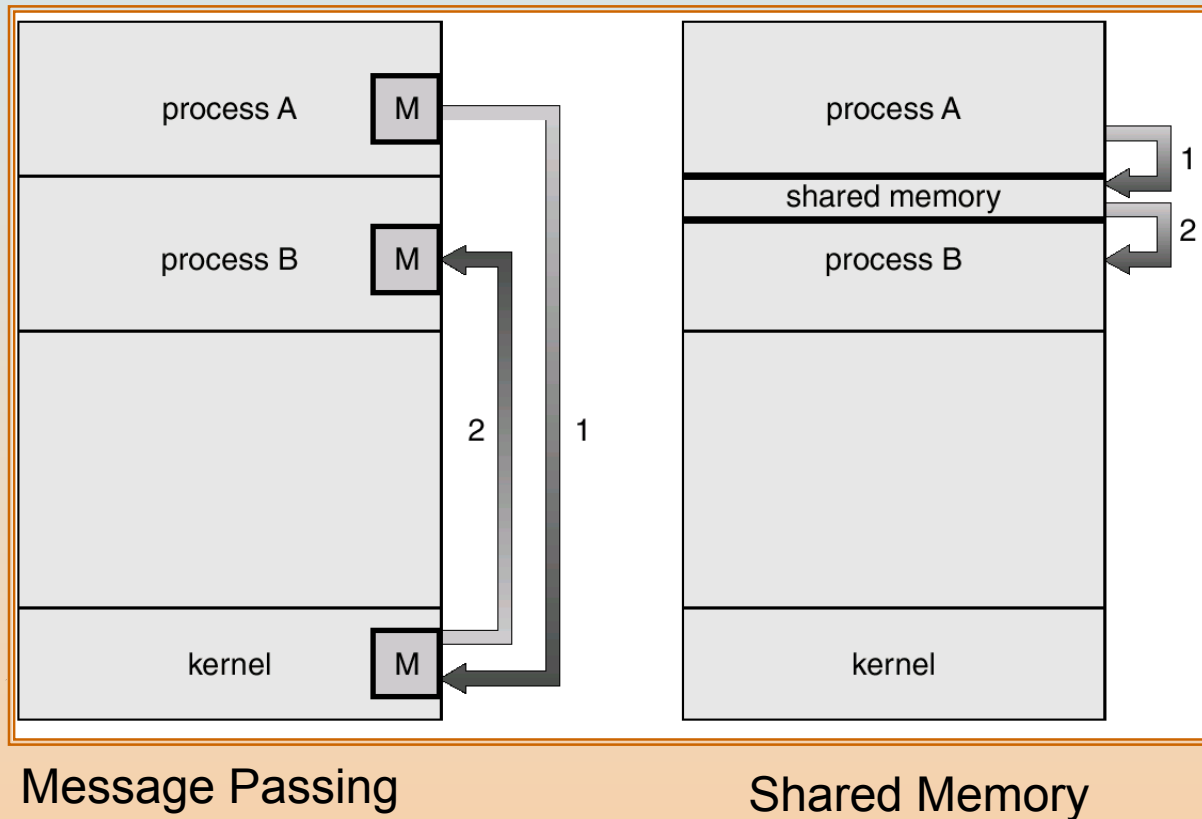
UNIX Running Multiple Programs





Communication Models

- Communication may take place using either message passing or shared memory





System Programs

- System programs provide a convenient environment for program development and execution. They can be divided into:
 - File manipulation
 - Status information
 - File modification
 - Programming language support
 - Program loading and execution
 - Communications
 - Application programs
- Most users' view of the operation system is defined by system programs, not the actual system calls





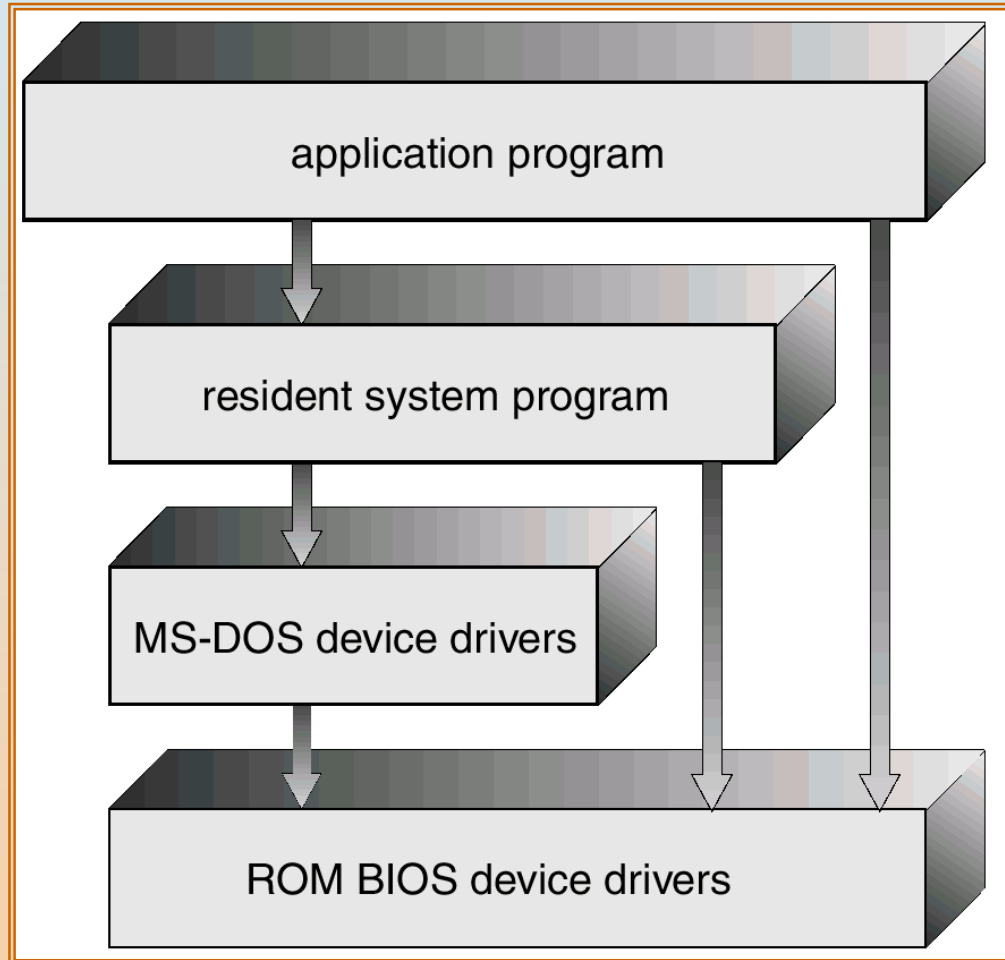
MS-DOS System Structure

- MS-DOS – written to provide the most functionality in the least space
 - Not divided into modules
 - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated





MS-DOS Layer Structure





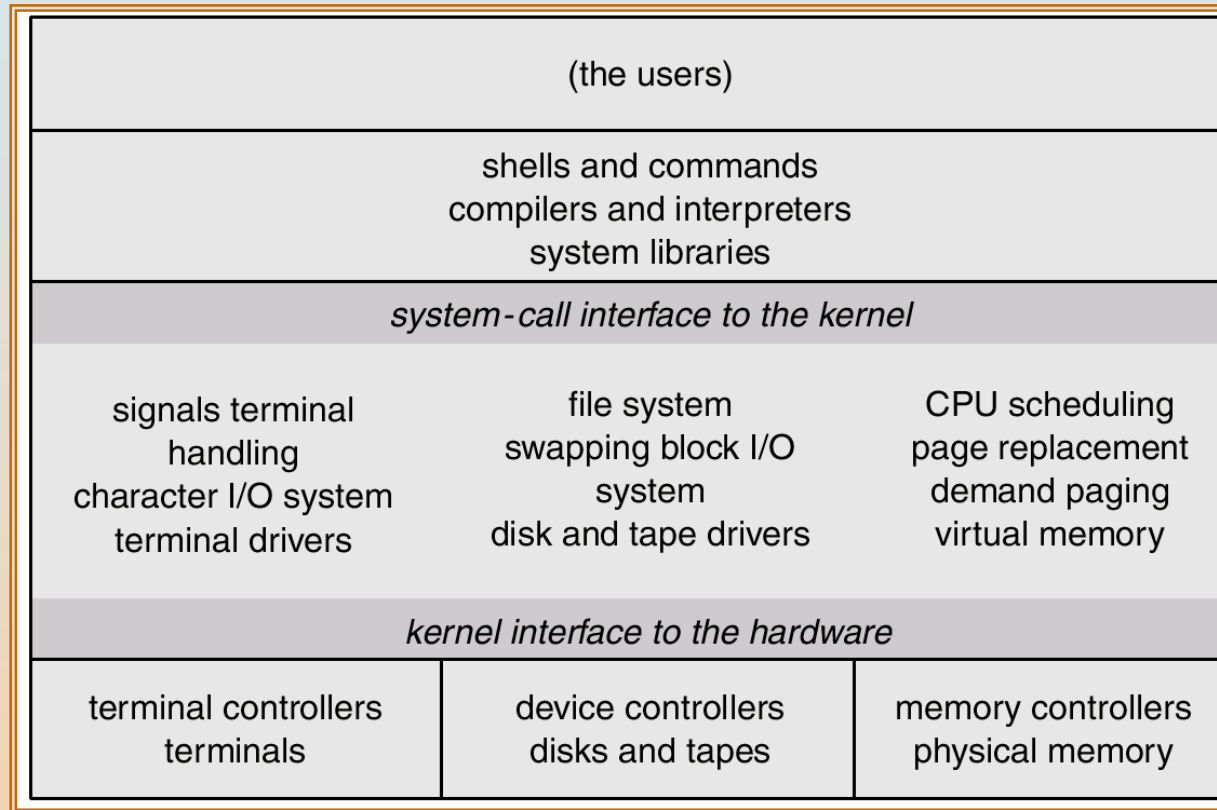
UNIX System Structure

- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring. The UNIX OS consists of two separable parts
 - Systems programs
 - The kernel
 - ▶ Consists of everything below the system-call interface and above the physical hardware
 - ▶ Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level





UNIX System Structure





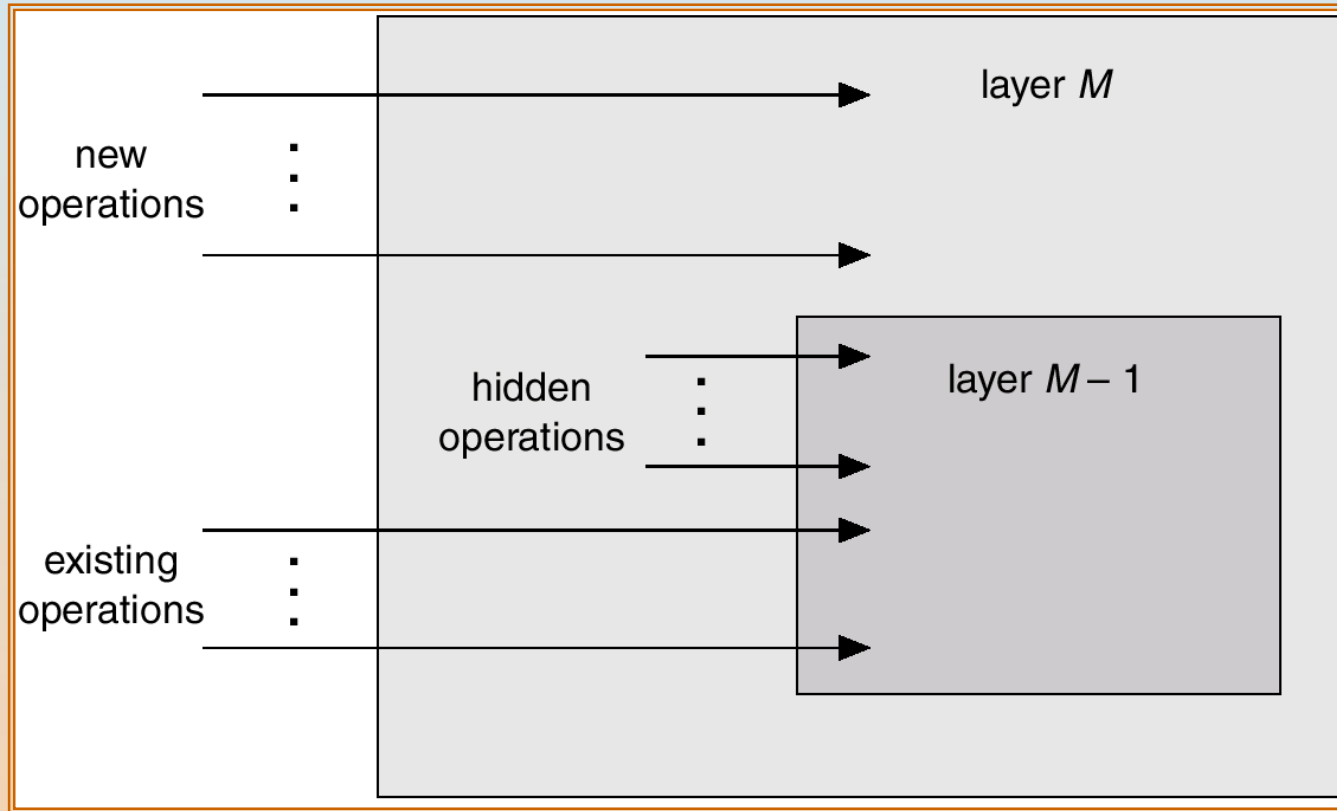
Layered Approach

- The operating system is divided into a number of layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers



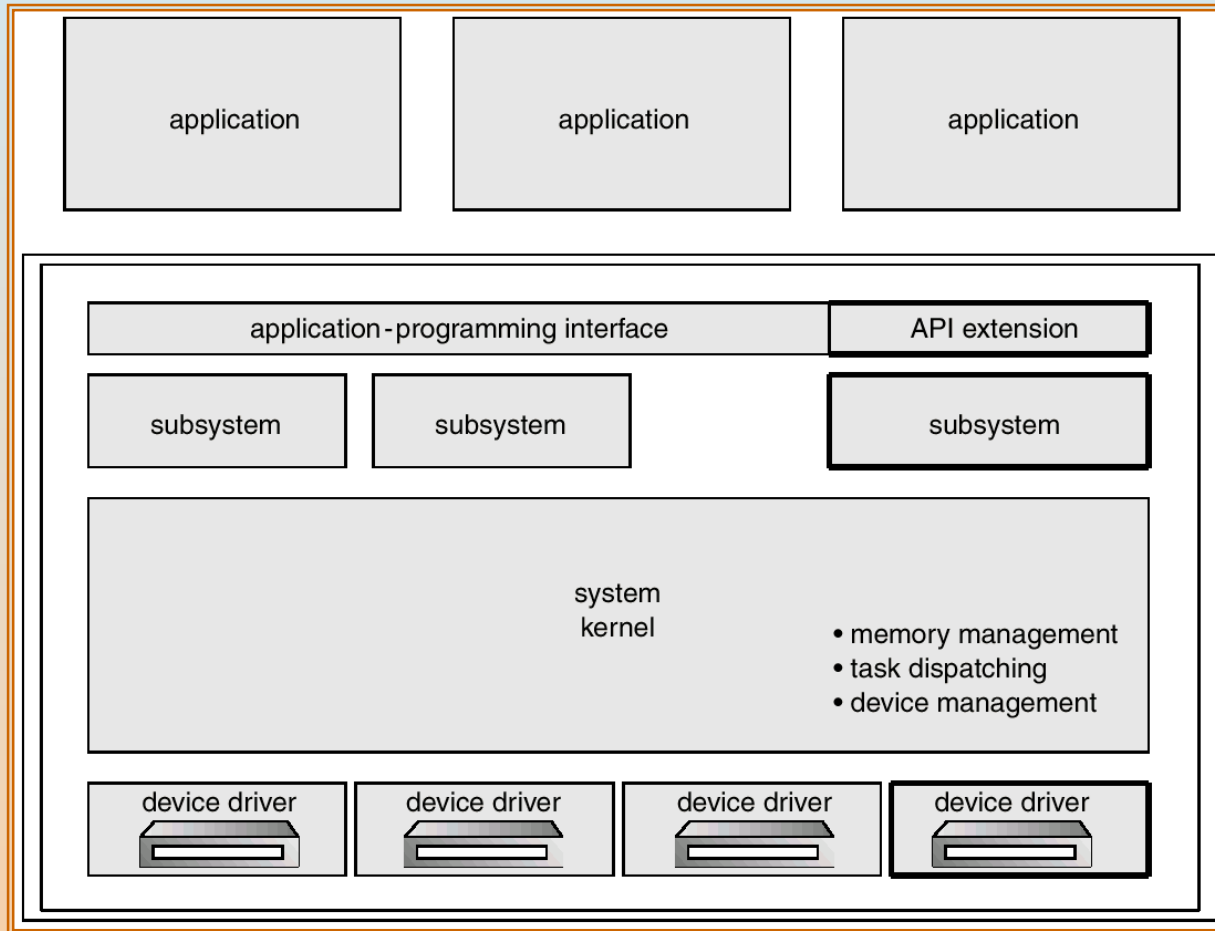


An Operating System Layer





OS/2 Layer Structure





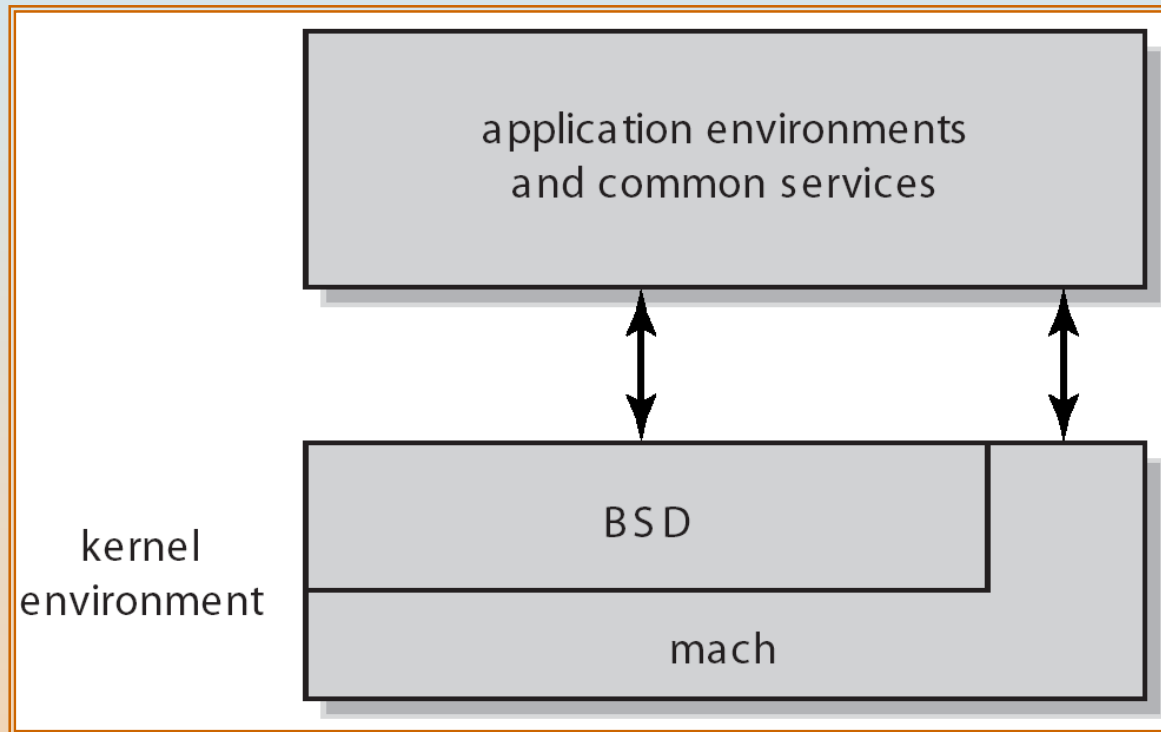
Microkernel System Structure

- Moves as much from the kernel into “*user*” space
- Communication takes place between user modules using message passing
- Benefits:
 - Easier to extend a microkernel
 - Easier to port the operating system to new architectures
 - More reliable (less code is running in kernel mode)
 - More secure
- Detriments:
 - Performance overhead of user space to kernel space communication



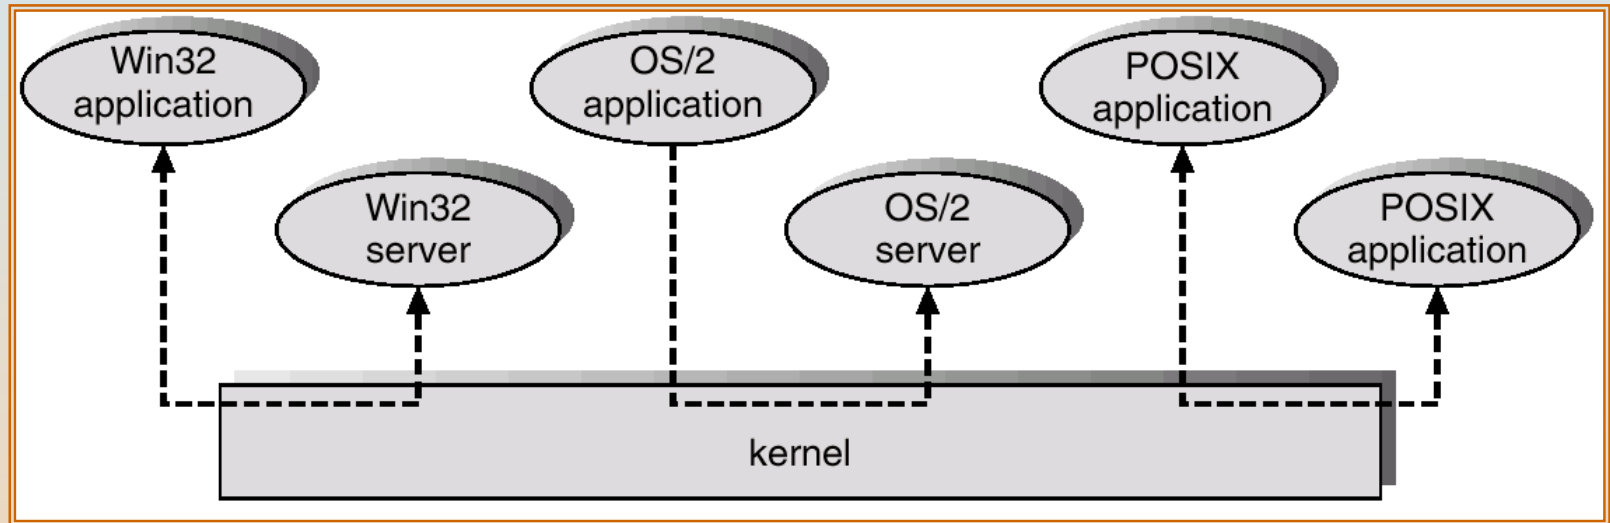


Mac OS X Structure





Windows NT Client-Server Structure





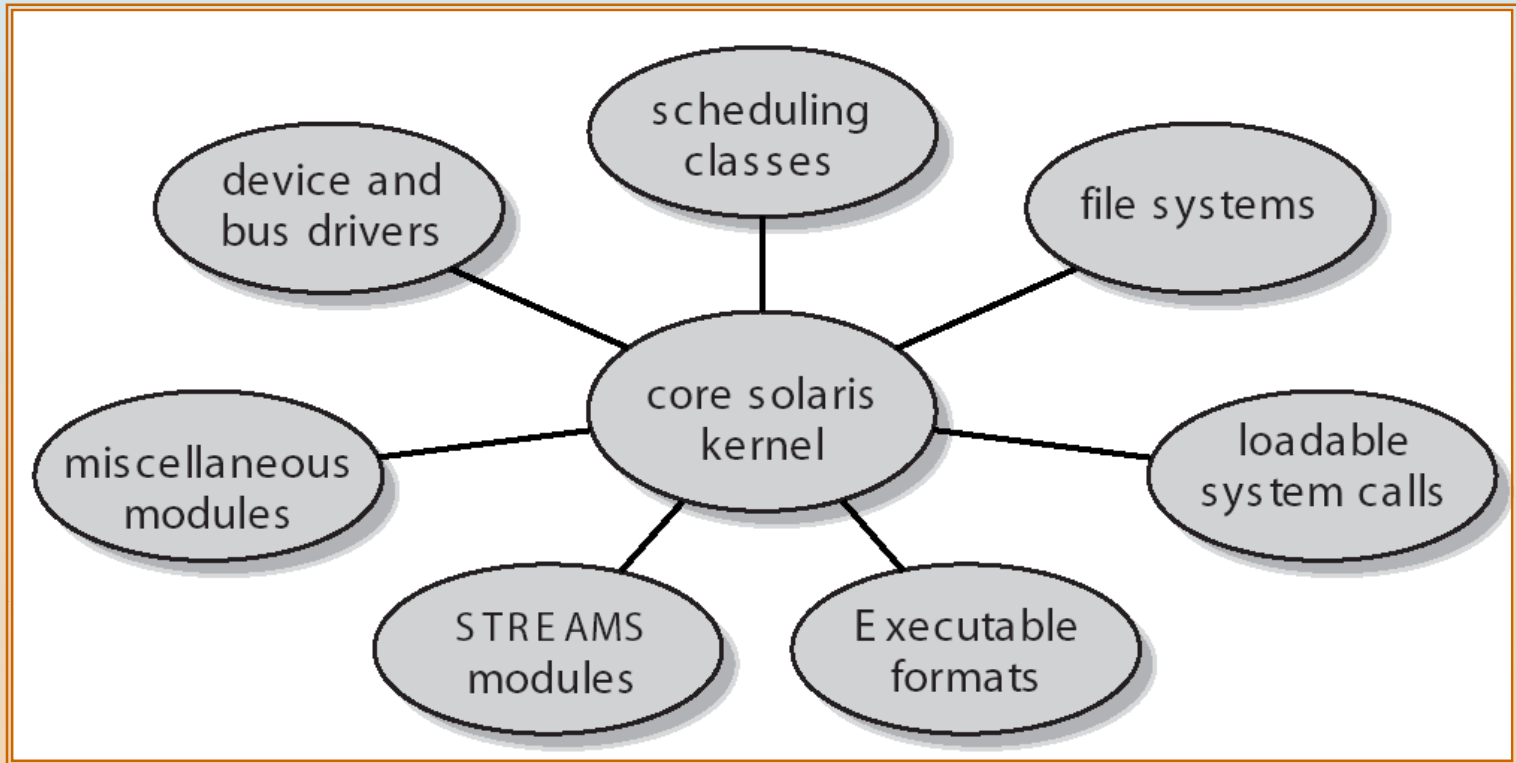
Modules

- Most modern operating systems implement kernel *modules*
 - Uses object-oriented approach
 - Each core component is separate
 - Each talks to the others over known interfaces
 - Each is loadable as needed within the kernel
- Overall, similar to layers but with more flexible





Solaris Modular Approach





Virtual Machines

- A *virtual machine* takes the layered approach to its logical conclusion. It treats hardware and the operating system kernel as though they were all hardware
- A virtual machine provides an interface *identical* to the underlying bare hardware
- The operating system creates the illusion of multiple processes, each executing on its own processor with its own (virtual) memory





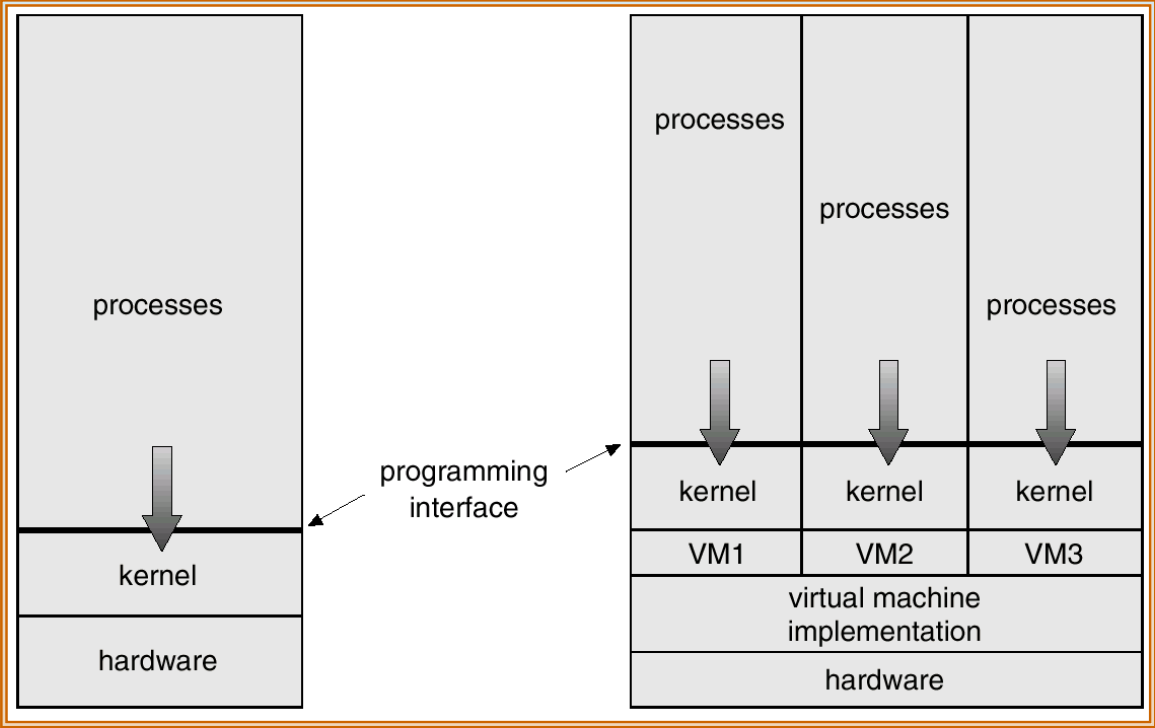
Virtual Machines (Cont.)

- The resources of the physical computer are shared to create the virtual machines
 - CPU scheduling can create the appearance that users have their own processor
 - Spooling and a file system can provide virtual card readers and virtual line printers
 - A normal user time-sharing terminal serves as the virtual machine operator's console





System Models



Non-virtual Machine

Virtual Machine





Advantages/Disadvantages of Virtual Machines

- The virtual-machine concept provides complete protection of system resources since each virtual machine is isolated from all other virtual machines. This isolation, however, permits no direct sharing of resources.
- A virtual-machine system is a perfect vehicle for operating-systems research and development. System development is done on the virtual machine, instead of on a physical machine and so does not disrupt normal system operation.
- The virtual machine concept is difficult to implement due to the effort required to provide an *exact* duplicate to the underlying machine






Java Virtual Machine

- Compiled Java programs are platform-neutral bytecodes executed by a Java Virtual Machine (JVM)
- JVM consists of
 - Class loader
 - Class verifier
 - Runtime interpreter
- Just-In-Time (JIT) compilers increase performance



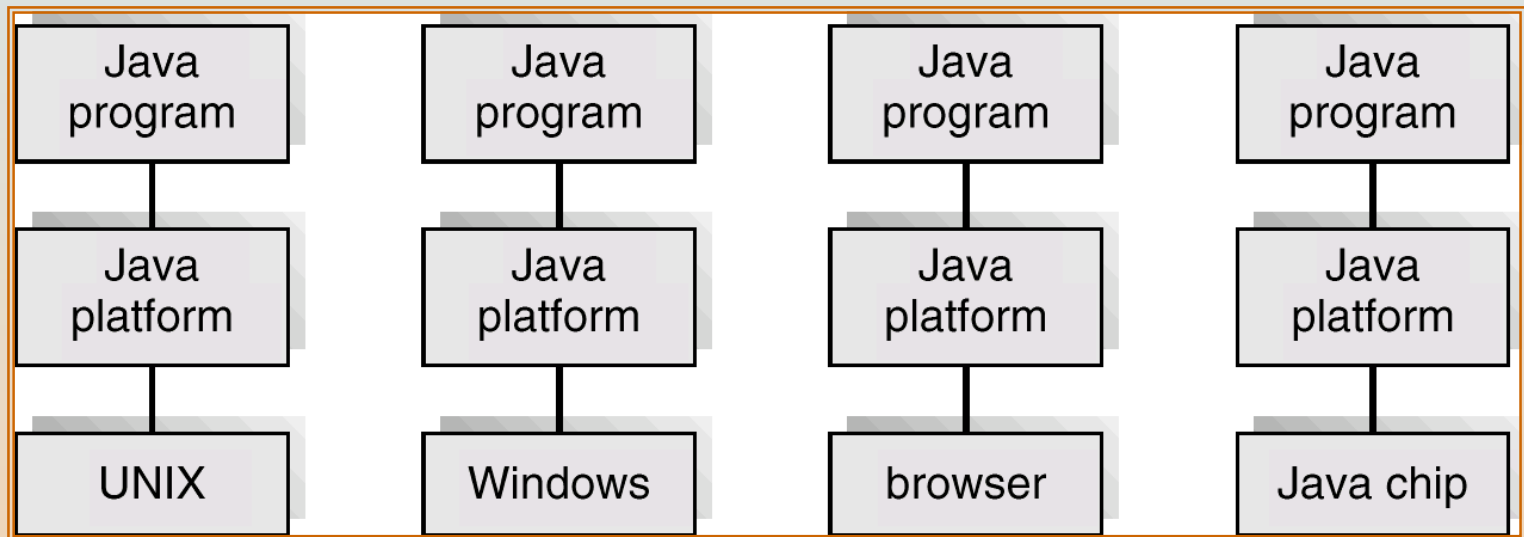


The Java Virtual Machine



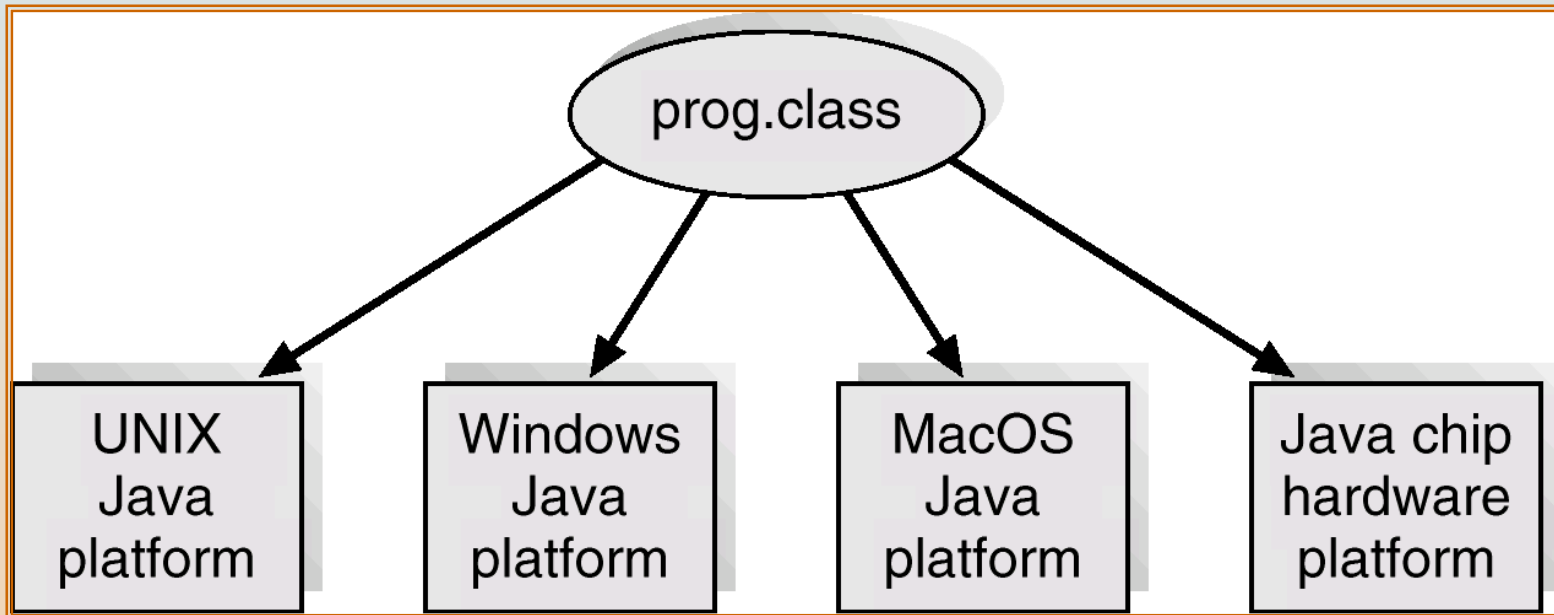


The Java Platform



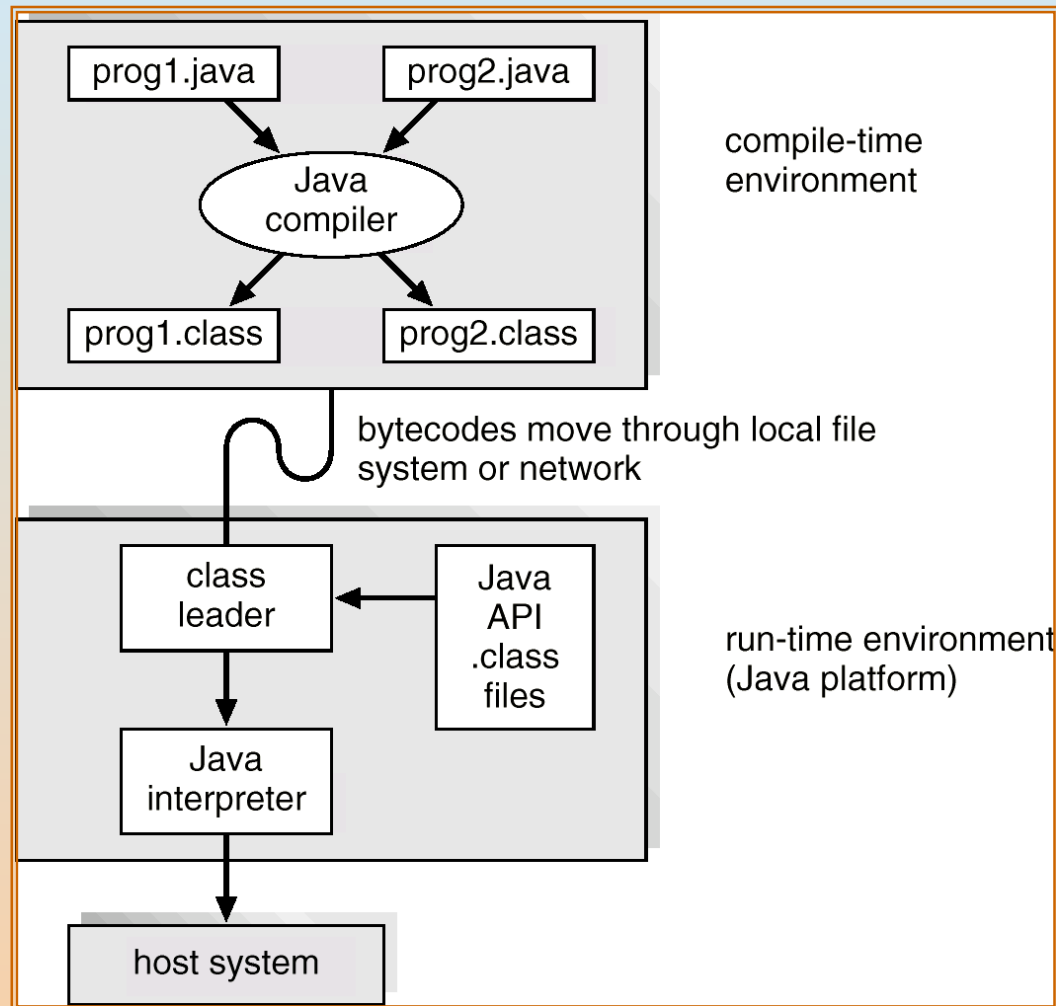


Java .class File on Cross Platforms





Java Development Environment





System Design Goals

- User goals – operating system should be convenient to use, easy to learn, reliable, safe, and fast
- System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient





Mechanisms and Policies

- Mechanisms determine how to do something, policies decide what will be done
- The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later





System Implementation

- Traditionally written in assembly language, operating systems can now be written in higher-level languages
- Code written in a high-level language:
 - Can be written faster
 - Is more compact.
 - Is easier to understand and debug
- An operating system is far easier to *port* (move to some other hardware) if it is written in a high-level language





System Design Goals

- User goals – operating system should be convenient to use, easy to learn, reliable, safe, and fast
- System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient





System Generation (SYSGEN)

- Operating systems are designed to run on any of a class of machines; the system must be configured for each specific computer site
- SYSGEN program obtains information concerning the specific configuration of the hardware system
- *Booting* – starting a computer by loading the kernel
- *Bootstrap program* – code stored in ROM that is able to locate the kernel, load it into memory, and start its execution

